



# Keras Models Module API Design Analysis

## 1. Constructor Required Parameters Tradeoff

**Required Parameter Counts:** The `Sequential` constructor requires **0** parameters (all constructor args have defaults, e.g. `layers=None`), whereas the `Functional` API model (using `Model(inputs, outputs)`) requires **2** parameters (`inputs` and `outputs` have no defaults). The base `Model` class itself can be constructed with **0** required args (it accepts `inputs=None` and `outputs=None` as optional) <sup>1</sup>. Thus, Functional requires two parameters vs. Sequential's zero (2 vs 0). In practice, this means a Functional model must be built with an explicit `Input` tensor and output tensor provided, unlike a Sequential model which infers its input from the first layer and needs no upfront specification <sup>2</sup>.

**Ratio (Functional:Sequential):** Functional's required parameter count is **2:0** relative to Sequential (Functional needs two mandatory arguments, Sequential needs none). In other words, Functional demands additional upfront information that Sequential does not. The specific extra parameter is the explicit `inputs` argument (and corresponding `outputs`) that Functional API models require, which Sequential does not. **Tradeoff:** Requiring an `inputs` tensor in Functional models yields greater flexibility (e.g. multiple or non-linear layer graphs, shared layers) at the cost of more complexity for the developer. The Sequential API's omission of an input requirement offers simplicity and convenience for linear stacks of layers, but at the expense of flexibility – Sequential cannot represent arbitrary layer graphs or multiple inputs/outputs without that extra specification <sup>2</sup>. This design choice illustrates the classic convenience vs. explicitness tradeoff: Functional makes the model architecture explicit (more parameters to specify, hence more verbose) in exchange for far greater modeling flexibility.

## 2. Serialization Naming Consistency

**Serialization Function Pairs:** Keras provides three pairs of functions for model serialization/deserialization. Two out of these three pairs follow a consistent `model_to_X` / `model_from_X` naming pattern, while one pair does not. Specifically:

- **JSON format:** `model.to_json()` and `keras.models.model_from_json()` – consistent naming (to/from JSON) <sup>3</sup>.
- **YAML format:** `model.to_yaml()` and `keras.models.model_from_yaml()` – consistent naming (to/from YAML).
- **Config dict format:** `model.get_config()` and `keras.Model.from_config()` – **inconsistent** (uses "get" vs "from", breaking the `to_X/from_X` symmetry) <sup>3</sup>.

**Consistency Ratio:** **2 out of 3** pairs have consistent naming, so the ratio is **2:3** (approximately 66% of the serialization pairs follow the pattern). The one format pair that breaks the pattern is the `config` serialization (using `get_config()` / `from_config()` instead of a `to_config()/from_config` convention). This violates the API design principle of *naming consistency*. The inconsistency exists largely for **legacy compatibility** – `get_config()` was an established naming convention in early Keras (for layers and

models) and has been retained for familiarity and backward-compatibility <sup>3</sup>. In other words, the designers favored not breaking the older API (`get_config`) even though it meant deviating from the newer `to_X/`  
`from_X` scheme, highlighting a consistency vs. legacy support tradeoff.

### 3. `Model.compile()` Signature Flexibility

**Default Parameter Counts:** In `Model.compile(...)`, there are several parameters with explicit non-`None` default values and several that default to `None`. Counting from the signature <sup>1</sup>: - **Explicit defaults (non-`None`):** 5 parameters - e.g. `optimizer="rmsprop"` (default optimizer), `run_eagerly=False`, `steps_per_execution=1`, `jit_compile="auto"`, and `auto_scale_loss=True` all have concrete default values <sup>1</sup>. - **Defaults to `None`:** 4 parameters - e.g. `loss=None`, `metrics=None`, `loss_weights=None`, `weighted_metrics=None` default to `None` (meaning the user typically should provide them) <sup>1</sup>.

This yields an **explicit:`None` default ratio of 5:4**. In other words, slightly more than half of the compile parameters have a built-in default, while the rest default to `None` (requiring user specification or using default behaviors). The explicit defaults are chosen for common convenience (e.g. a default optimizer and sensible booleans), whereas `None` defaults indicate that those arguments (like loss or metrics) usually need to be explicitly defined by the developer.

**Most Permissive Parameter:** The `metrics` argument of `compile()` has the most permissive type signature. It accepts a wide range of types: a string (metric name), a function, a `tf.keras.metrics.Metric` object, or a list/dict of these for multi-output models <sup>4</sup>. (Similarly, the `loss` argument is also very flexible, allowing strings, loss objects, callables, and lists/dicts for multiple outputs, but `metrics` is especially diverse in accepted forms.) This high permissiveness reflects an API philosophy favoring **flexibility and developer convenience** over strict type safety. The user can provide metrics in whatever form is easiest (built-in names, custom functions, etc.), and Keras will interpret it appropriately <sup>4</sup>. The tradeoff is that this implicit type-handling reduces static type safety – the API must internally check and convert types at runtime. In short, Keras opts to be forgiving and flexible in `compile()` inputs (to streamline common use cases) at the cost of potential ambiguity or error catching only at runtime, exemplifying the flexibility vs. type-safety design balance.

### 4. Inheritance Hierarchy Complexity

**Model Inheritance:** The immediate parent class of `Model` is the base Keras `Layer` class (all Keras models are a specialized form of Layer). `Model` uses single inheritance – it inherits from one parent class (`Layer`) and does not use multiple inheritance.

**Inheritance Depth:** - `Model` has an inheritance depth of **2** levels from Python's built-in `object` (`object → Layer → Model`). - **Functional** models in Keras (internally often represented by the `Functional` subclass) inherit from `Model`, so their depth is **3** (`object → Layer → Model → Functional`). - **Sequential** is also a subclass of `Model` <sup>5</sup>, giving it a depth of **3** as well (`object → Layer → Model → Sequential`).

Thus, the more “complex” convenience classes like Functional/Sequential add one extra layer of inheritance on top of the base Model. The difference in depth between the simplest class (base `Model`, depth 2) and the most layered subclass (e.g. `Functional` or `Sequential`, depth 3) is **1 level**.

**Tradeoff (Reuse vs. Complexity):** This hierarchy reveals a tradeoff between **code reuse and architectural complexity**. By subclassing `Model`, both Sequential and Functional APIs inherit all of Model’s functionality (training loops, serialization, etc.) without reimplementing it <sup>5</sup>. This is a clear code-reuse benefit. However, adding inheritance levels increases the complexity of the class hierarchy (and the learning curve for understanding the framework’s internals). Keras accepted a modest increase in complexity (one extra inheritance level) to maximize reuse of the base Model logic. The result is that `Sequential` and `Functional` are thin wrappers that leverage `Model`’s features rather than duplicating code, illustrating the design preference for reuse and maintainability even if it means a deeper inheritance tree.

## 5. Comparison Table of Keras Model Classes

The table below compares key aspects of the `Functional`, `Model`, and `Sequential` classes. It has **5 columns – Class Name, Required Parameters, Optional Parameters, Inheritance Depth, and Primary Use Case** – and **3 rows** (one for each class), sorted alphabetically by class name. “Required Parameters” is the count of constructor parameters without default values, “Optional Parameters” is the count of constructor parameters that have default values, “Inheritance Depth” is the number of inheritance levels from the base `object` (i.e. how many classes in the MRO chain above it), and “Primary Use Case” categorizes the class’s typical role in Keras.

Class Name	Required Parameters	Optional Parameters	Inheritance Depth	Primary Use Case
<code>Functional</code>	2 (e.g. <code>inputs</code> , <code>outputs</code> )	2 (e.g. <code>name</code> , <code>trainable</code> )	3 levels from <code>object</code>	Graph Models
<code>Model</code>	0 (none required)	3 (e.g. <code>inputs=None</code> , <code>outputs=None</code> , <code>name=None</code> )	2 levels from <code>object</code>	Base Abstraction
<code>Sequential</code>	0 (none required)	3 (e.g. <code>layers=None</code> , <code>trainable=True</code> , <code>name=None</code> )	3 levels from <code>object</code>	Linear Models

Each class builds on the prior abstractions: `Model` is the foundational class (providing core training/evaluation functionality), `Functional` (often represented via `keras.Model(inputs, outputs)`) extends Model for arbitrary network graphs, and `Sequential` extends Model for the specific case of linear layer stacks. The counts and depths reflect how simplicity vs. flexibility is managed – Sequential keeps constructor requirements minimal and inherits one extra layer for convenience, while Functional requires more inputs for flexibility, both leveraging the base Model’s capabilities for code reuse <sup>5</sup>.

**Sources:** Key insights and data were gathered from the TensorFlow Keras source and documentation <sup>1</sup> <sup>4</sup> <sup>3</sup> <sup>2</sup> <sup>5</sup>, which illustrate the design decisions in the Keras models module as of Feb 2026. Each design choice (required constructor args, naming schemes, compile signature, inheritance structure)

highlights a balance between developer-friendly simplicity and powerful flexibility in the API. The analysis above quantifies these choices and explains the tradeoffs involved.

---

1 4 Model training APIs

[https://keras.io/api/models/model\\_training\\_apis/](https://keras.io/api/models/model_training_apis/)

2 5 python - Conversion of a sequential model to a functional model with Keras 2.2.0 - Stack Overflow

<https://stackoverflow.com/questions/50837715/conversion-of-a-sequential-model-to-a-functional-model-with-keras-2-2-0>

3 Save, serialize, and export models

[https://keras.io/guides/serialization\\_and\\_saving/](https://keras.io/guides/serialization_and_saving/)